

# Neural networks

COMS 4771 Fall 2019

## Overview

- ▶ Structure and power of neural networks
- ▶ Backpropagation
- ▶ Practical issues
- ▶ Convolutions

0 / 26

1 / 26

## Parametric featurizations I

- ▶ So far: features ( $x$  or  $\varphi(x)$ ) are fixed during training
  - ▶ Consider a (small) collection of feature transformations  $\varphi$
  - ▶ Select  $\varphi$  via cross-validation – outside of normal training
- ▶ “Deep learning” approach:
  - ▶ Use  $\varphi$  with many tunable parameters
  - ▶ Optimize parameters of  $\varphi$  during normal training process

2 / 26

## Parametric featurizations II

- ▶ Neural network: parameterization for function  $f: \mathbb{R}^d \rightarrow \mathbb{R}$ 
  - ▶  $f(x) = \varphi(x)^T w$
  - ▶ Parameters include both  $w$  and parameters of  $\varphi$
  - ▶ Varying parameters of  $\varphi$  allows  $f$  to be essentially any function!
  - ▶ Major challenge: optimization (a lot of tricks to make it work)




Figure 1: Neural network

3 / 26

# Feedforward neural network

- Architecture of a feedforward neural network
  - Directed acyclic graph  $G = (V, E)$
  - One source node (vertex) per input, one sink node per output
  - Other nodes are hidden units
  - Each edge  $(u, v) \in E$  has a weight parameter  $w_{u,v} \in \mathbb{R}$
  - Value  $h_v$  of node  $v$  given values of parents is

$$h_v := \sigma_v(z_v), \quad z_v := \sum_{u \in V: (u,v) \in E} w_{u,v} \cdot h_u.$$

- $\sigma_v: \mathbb{R} \rightarrow \mathbb{R}$  is the activation function (e.g., sigmoid)




Figure 2: Feedforward neural network architecture

4 / 26

# Standard layered architectures

- Standard architecture arranges nodes into sequence of  $L$  layers
  - Edges only go from one layer to the next
  - Can write function using matrices of weight parameters
- $$f(\mathbf{x}) = \sigma_L(\mathbf{W}_L \sigma_{L-1}(\cdots \sigma_1(\mathbf{W}_1 \mathbf{x}) \cdots))$$
- $d_\ell$  nodes in layer  $\ell$ ;  $\mathbf{W}_\ell \in \mathbb{R}^{d_\ell \times d_{\ell-1}}$  are weight parameters
  - Activation function  $\sigma_\ell: \mathbb{R} \rightarrow \mathbb{R}$  is applied coordinate-wise to input
- Often also include “bias” parameters  $\mathbf{b}_\ell \in \mathbb{R}^{d_\ell}$

$$f(\mathbf{x}) = \sigma_L(\mathbf{b}_L + \mathbf{W}_L \sigma_{L-1}(\cdots \sigma_1(\mathbf{b}_1 + \mathbf{W}_1 \mathbf{x}) \cdots))$$

- Tunable parameters:  $\boldsymbol{\theta} = (\mathbf{W}_1, \mathbf{b}_1, \dots, \mathbf{W}_L, \mathbf{b}_L)$




Figure 3: Standard feedforward architecture

5 / 26

# Well-known activation functions

- Heaviside:  $\sigma(z) = \mathbb{1}_{\{z \geq 0\}}$ 
  - Popular in the 1940s; also called step function
- Sigmoid (from logistic regression):  $\sigma(z) = 1/(1 + e^{-z})$ 
  - Popular since 1970s
- Hyperbolic tangent:  $\sigma(z) = \tanh(z)$ 
  - Similar to sigmoid, but range is  $(-1, 1)$  rather than  $(0, 1)$
- Rectified Linear Unit (ReLU):  $\sigma(z) = \max\{0, z\}$ 
  - Popular since 2012
- Identity:  $\sigma(z) = z$ 
  - Popular with luddites
- Softmax:  $\sigma(\mathbf{v})_i = \exp(v_i) / \sum_j \exp(v_j)$ 
  - Special vector-valued activation function
  - Popular for final layer in multi-class classification

6 / 26

# Power of non-linear activations

- What happens if every activation function is linear/affine?

7 / 26

# Necessity of multiple layers

- Suppose only have input and output layers, so function  $f$  is
$$f(\mathbf{x}) = \sigma(b + \mathbf{w}^\top \mathbf{x})$$
where  $b \in \mathbb{R}$  and  $\mathbf{w} \in \mathbb{R}^d$  (so  $\mathbf{w}^\top \in \mathbb{R}^{1 \times d}$ )
- If  $\sigma$  is monotone (e.g., Heaviside, sigmoid, hyperbolic tangent, ReLU, identity), then  $f$  has same limitations as a linear/affine classifier




Figure 4: XOR problem

# Neural network approximation theorems

- Theorem** (Cybenko, 1989; Hornik, Stinchcombe, & White, 1989): For any continuous function  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  and any  $\varepsilon > 0$ , there is a two-layer neural network (with parameters  $\theta = (\mathbf{W}_1, \mathbf{b}_1, \mathbf{w}_2)$ ) s.t.

$$\max_{\mathbf{x} \in [0,1]^d} |f(\mathbf{x}) - \mathbf{w}_2^\top \sigma_1(\mathbf{b}_1 + \mathbf{W}_1 \mathbf{x})| < \varepsilon.$$

Here,  $\sigma_1$  can be any non-linear activation function from above.

- Many caveats
  - “Width” (number of hidden units) may need to be very large
  - Does not tell us how to find the network
  - Does not justify deeper networks

# Fitting neural networks to data

- Training data  $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n) \in \mathbb{R}^d \times \mathcal{Y}$
- Fix architecture:  $G = (V, E)$  and activation functions
- Plug-in principle: find parameters  $\theta$  of neural network  $f_\theta$  to minimize empirical risk (possibly with a surrogate loss)

$$\begin{aligned} \widehat{\mathcal{R}}(\theta) &= \frac{1}{n} \sum_{i=1}^n (f_\theta(\mathbf{x}_i) - y_i)^2 && \text{regression} \\ \widehat{\mathcal{R}}(\theta) &= \frac{1}{n} \sum_{i=1}^n \ell_{\log}(-y_i f_\theta(\mathbf{x}_i)) && \text{binary classification} \\ \widehat{\mathcal{R}}(\theta) &= \frac{1}{n} \sum_{i=1}^n \ell_{\text{ce}}(\tilde{\mathbf{y}}_i, f_\theta(\mathbf{x}_i)) && \text{multi-class classification} \end{aligned}$$

- (Could use other surrogate loss functions ...)
- Typically objective is not convex in parameters  $\theta$
  - Nevertheless, local search (e.g., SGD) often works well!

# Backpropagation

- Backpropagation (backprop)**: Algorithm for computing partial derivatives wrt weights in a feedforward neural network
  - Clever organization of partial derivative computations with [chain rule](#)
  - Use in combination with gradient descent, SGD, etc.
- Consider loss on a single example  $(\mathbf{x}, y)$ , written  $J := \ell(y, f_\theta(\mathbf{x}))$
- Goal: compute  $\frac{\partial J}{\partial w_{u,v}}$  for every edge  $(u, v) \in E$
- Initial step of backprop: **forward propagation**
  - Compute  $z_v$ 's and  $h_v$ 's for every node  $v \in V$
  - Running time: linear in size of network
- Rest of backprop also just requires time linear in size of network!

Derivative of loss with respect to weights

- ▶ Let  $\hat{y}_1, \hat{y}_2, \dots$  denote the values at the output nodes.
- ▶ Then by chain rule,


$$\frac{\partial J}{\partial w_{u,v}} = \sum_i \frac{\partial J}{\partial \hat{y}_i} \cdot \frac{\partial \hat{y}_i}{\partial w_{u,v}}.$$

Derivative of output with respect to weights

- ▶ Assume for simplicity there is just a single output,  $\hat{y}$
- ▶ Chain rule, again:

$$\frac{\partial \hat{y}}{\partial w_{u,v}} = \frac{\partial \hat{y}}{\partial h_v} \cdot \frac{\partial h_v}{\partial w_{u,v}}$$


- ▶ First term: trickier; we'll handle later
- ▶ Second term:



Derivative of output with respect to hidden units

- ▶ Compute  $\frac{\partial \hat{y}}{\partial h_v}$  for all vertices in decreasing order of layer number
- ▶ If  $v$  is not the output node, then by chain rule (yet again),

$$\frac{\partial \hat{y}}{\partial h_v} = \sum_{v': (v,v') \in E} \frac{\partial \hat{y}}{\partial h_{v'}} \cdot \frac{\partial h_{v'}}{\partial h_v}$$



Example: chain graph I




Figure 5: Chain graph; assume same activation  $\sigma$  in every layer

- ▶ Parameters  $\theta = (w_{0,1}, w_{1,2}, \dots, w_{L-1,L})$
- ▶ Fix input value  $x \in \mathbb{R}$ ; what is  $\frac{\partial h_L}{\partial w_{i-1,i}}$  for  $i = 1, \dots, L$ ?
- ▶ Forward propagation:
  - ▶  $h_0 := x$
  - ▶ For  $i = 1, 2, \dots, L$ :

$$z_i := w_{i-1,i} h_{i-1}$$

$$h_i := \sigma(z_i)$$

# Example: chain graph II




Figure 6: Chain graph; assume same activation  $\sigma$  in every layer

# Practical issues I: Initialization

- ▶ Backprop:
  - ▶ For  $i = L, L - 1, \dots, 1$ :

$$\frac{\partial h_L}{\partial h_i} := \begin{cases} 1 & \text{if } i = L \\ \frac{\partial h_L}{\partial h_{i+1}} \cdot \sigma'(z_{i+1}) w_{i,i+1} & \text{if } i < L \end{cases}$$
$$\frac{\partial h_L}{\partial w_{i-1,i}} := \frac{\partial h_L}{\partial h_i} \cdot \sigma'(z_i) h_{i-1}$$

- ▶ Ensure inputs are standardized: every feature has zero mean and unit variance (wrt training data)
- ▶ Initialize weights randomly for gradient descent / SGD

# Practical issues II: Architecture choice

# Convolutional nets

- ▶ Architecture can be regarded as a “hyperparameter”
- ▶ Optimization-inspired architecture choice
  - ▶ With wide enough network, can get training error rate zero
  - ▶ Use the smallest network that lets you get zero training error rate
  - ▶ Then add regularization term to objective (e.g., sum of squares of weights), and optimize the regularized ERM objective
- ▶ Entire research communities are trying to figure out good architectures for their problems

- ▶ Neural networks with convolutional layers
  - ▶ Useful when inputs have locality structure
  - ▶ Sequential structure (e.g., speech waveform)
  - ▶ 2D grid structure (e.g., image)
  - ▶ ...
- ▶ Weight matrix  $\mathbf{W}_\ell$  is highly-structured
  - ▶ Determined by a small filter
  - ▶ Time to compute  $\mathbf{W}_\ell \mathbf{h}_{\ell-1}$  is typically  $\ll d_\ell \times d_{\ell-1}$  (e.g., closer to  $\max\{d_\ell, d_{\ell-1}\}$ )

# Convolutions I

- Convolution of two continuous functions:  $h := f * g$

$$h(x) = \int_{-\infty}^{+\infty} f(y)g(x - y) \, dy$$

- If  $f(x) = 0$  for  $x \notin [-w, +w]$ , then

$$h(x) = \int_{-w}^{+w} f(y)g(x - y) \, dy$$

- Replaces  $g(x)$  with weighted combination of  $g$  at nearby points
- For functions on discrete domain, replace integral with sum

$$h(i) = \sum_{j=-\infty}^{\infty} f(j)g(i - j)$$

# Convolutions II

- E.g., suppose only  $f(0), f(1), f(2)$  are non-zero, and  $g$  is zero-padded (in this case,  $g(i) = 0$  for  $i < 1$  or  $i > 5$ ). Then:

$$\begin{bmatrix} h(1) \\ h(2) \\ h(3) \\ h(4) \\ h(5) \\ h(6) \\ h(7) \end{bmatrix} = \begin{bmatrix} f(0) & 0 & 0 & 0 & 0 \\ f(1) & f(0) & 0 & 0 & 0 \\ f(2) & f(1) & f(0) & 0 & 0 \\ 0 & f(2) & f(1) & f(0) & 0 \\ 0 & 0 & f(2) & f(1) & f(0) \\ 0 & 0 & 0 & f(2) & f(1) \\ 0 & 0 & 0 & 0 & f(2) \end{bmatrix} \begin{bmatrix} g(1) \\ g(2) \\ g(3) \\ g(4) \\ g(5) \end{bmatrix}$$




Figure 7: Convolutional layer

# 2D convolutions I

- Similar for 2D inputs (e.g., images), except now sum over two indices
  - $g$  is the input image
  - $f$  is the filter
  - Lots of variations (e.g., padding, strides, multiple “channels”)




Figure 8: 2D convolution, with padding, no stride

# 2D convolutions II

- Similar for 2D inputs (e.g., images), except now sum over two indices
  - $g$  is the input image
  - $f$  is the filter
  - Lots of variations (e.g., padding, strides, multiple “channels”)




Figure 9: 2D convolution, with padding, no stride

## 2D convolutions III

- ▶ Similar for 2D inputs (e.g., images), except now sum over two indices
  - ▶  $g$  is the input image
  - ▶  $f$  is the filter
  - ▶ Lots of variations (e.g., padding, strides, multiple “channels”)




Figure 10: 2D convolution, with padding, no stride

24 / 26

## 2D convolutions IV

- ▶ Similar for 2D inputs (e.g., images), except now sum over two indices
  - ▶  $g$  is the input image
  - ▶  $f$  is the filter
  - ▶ Lots of variations (e.g., padding, strides, multiple “channels”)




Figure 11: 2D convolution, with padding, no stride

- ▶ Use additional layers/activations to down-sample after convolution

25 / 26

## Postscript: Tangent model

- ▶ Let  $f_{\theta}: \mathbb{R}^d \rightarrow \mathbb{R}$  be a neural network function with parameters  $\theta \in \mathbb{R}^p$  ( $p$  is total number of parameters)
- ▶ Fix  $x$ , and consider first-order approximation of  $f_{\theta}(x)$  around  $\theta = \theta^{(0)}$ :

$$f_{\theta}(x) \approx f_{\theta^{(0)}}(x) + \nabla f_{\theta^{(0)}}(x)^{\top}(\theta - \theta^{(0)})$$

Here,  $\nabla$  is gradient wrt parameters  $\theta$ , not wrt input  $x$

- ▶ Consider feature transformation  $\varphi(x) := \nabla f_{\theta^{(0)}}(x)$ , determined entirely by initial parameters  $\theta^{(0)}$
- ▶ If the first-order approximation is accurate (i.e., we never consider  $\theta$  too far from  $\theta^{(0)}$ ), then back to a linear model (over  $p$  features  $\varphi(x) \in \mathbb{R}^p$ )
  - ▶ Called the [\*tangent model\*](#)
- ▶ Upshot: To really exploit power of “deep learning”, we must be changing parameters a lot during training!

26 / 26